

```
In [1]: '''  
38-Cloud: Cloud Segmentation in Satellite Images  
'''
```

```
Out[1]: '\n38-Cloud: Cloud Segmentation in Satellite Images\n'
```

```
In [2]: import os  
os.environ['KMP_DUPLICATE_LIB_OK'] = 'TRUE'  
  
import numpy as np  
from PIL import Image  
from torch.utils.data import Dataset  
import torch  
import torch.nn as nn  
import torch.nn.functional as F  
from pathlib import Path  
import matplotlib.pyplot as plt  
from torch.utils.data import DataLoader, random_split
```



```

In [3]: class CustomCloudDataset(Dataset):
    def __init__(self, red_dir, green_dir, blue_dir, nir_dir, gt_dir, transform=None):
        self.red_files = list(Path(red_dir).glob('*.*tif'))
        self.green_dir = green_dir
        self.blue_dir = blue_dir
        self.nir_dir = nir_dir
        self.gt_dir = gt_dir
        self.transform = transform

    def __len__(self):
        return len(self.red_files)

    def __getitem__(self, idx):
        file_name = self.red_files[idx].name

        red = Image.open(self.red_files[idx])
        green = Image.open(Path(self.green_dir) / file_name.replace('red', 'green'))
        blue = Image.open(Path(self.blue_dir) / file_name.replace('red', 'blue'))
        nir = Image.open(Path(self.nir_dir) / file_name.replace('red', 'nir'))
        gt = Image.open(Path(self.gt_dir) / file_name.replace('red', 'gt'))

        image = np.stack([np.array(red), np.array(green), np.array(blue), np.array(nir)], axis=2).astype(np.float32)
        mask = np.array(gt).astype(np.uint8)
        mask = np.where(mask == 255, 1, 0).astype(np.uint8)

        if self.transform:
            image = self.transform(image)
            mask = self.transform(mask)

        image = torch.from_numpy(image).float().permute(2, 0, 1)
        mask = torch.from_numpy(mask).long()

        return image, mask

    def open_as_array(self, idx, include_nir=True):
        file_name = self.red_files[idx].name

        red = np.array(Image.open(self.red_files[idx]))
        green = np.array(Image.open(Path(self.green_dir) / file_name.replace('red', 'green')))
        blue = np.array(Image.open(Path(self.blue_dir) / file_name.replace('red', 'blue')))

```

```

    if include_nir:
        nir = np.array(Image.open(Path(self.nir_dir) / file_name.replace('red', 'nir')))
        image = np.stack([red, green, blue, nir], axis=2)
    else:
        image = np.stack([red, green, blue], axis=2)

    # Normalize the image
    max_val = np.iinfo(image.dtype).max
    image = image.astype(np.float32) / max_val

    return image

def open_mask(self, idx):
    file_name = self.red_files[idx].name
    mask = np.array(Image.open(Path(self.gt_dir) / file_name.replace('red', 'gt')))
    mask = np.where(mask == 255, 1, 0)
    return mask

```

```

In [4]: path = "./38-Cloud_training/"
        red_dir = path + "train_red"
        green_dir = path + "train_green"
        blue_dir = path + "train_blue"
        nir_dir = path + "train_nir"
        gt_dir = path + "train_gt"

        dataset = CustomCloudDataset(red_dir, green_dir, blue_dir, nir_dir, gt_dir)
        print(len(dataset))

```

8400



```

In [5]: x, y = dataset[1200]
        x.shape, y.shape

```

Out[5]: (torch.Size([4, 384, 384]), torch.Size([384, 384]))

```
In [6]: '''  
The ideas for the drawing portion were inspired by this article:  
https://www.kaggle.com/code/cordmaur/38-cloud-data-preparation?scriptVersionId=31357028  
'''  
  
fig, ax = plt.subplots(1, 2, figsize=(10, 9))  
ax[0].imshow(dataset.open_as_array(1000, include_nir=False))  
ax[1].imshow(dataset.open_mask(1000), cmap='gray')  
  
plt.show()
```



```
In [7]: total_size = len(dataset)
train_size = int(0.7 * total_size)
val_size = int(0.15 * total_size)
#train_size = int(0.1 * total_size)
#val_size = int(0.01 * total_size)
test_size = total_size - train_size - val_size

# Split the dataset
train_dataset, val_dataset, test_dataset = random_split(dataset, [train_size, val_size, test_size])

# Create DataLoaders for each set
batch_size = 4
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)

# Check if GPU is available and set the device accordingly
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
#device = torch.device("cpu")

print(f"Using device: {device}")
```

Using device: cuda


```

In [8]: class ConvBlock(nn.Module):
        def __init__(self, in_channels, out_channels):
            super(ConvBlock, self).__init__()
            self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3, padding=1)
            self.bn1 = nn.BatchNorm2d(out_channels)
            self.relu = nn.ReLU(inplace=True)
            self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size=3, padding=1)
            self.bn2 = nn.BatchNorm2d(out_channels)

        def forward(self, x):
            x = self.relu(self.bn1(self.conv1(x)))
            x = self.relu(self.bn2(self.conv2(x)))
            return x

class UpBlock(nn.Module):
    def __init__(self, in_channels, middle_channels, out_channels):
        super(UpBlock, self).__init__()
        self.up = nn.ConvTranspose2d(in_channels, middle_channels, kernel_size=2, stride=2)
        self.conv_block = ConvBlock(middle_channels + out_channels, out_channels)

    def forward(self, x, skip):
        x = self.up(x)
        x = torch.cat([x, skip], dim=1)
        return self.conv_block(x)

class UNet(nn.Module):
    def __init__(self, in_channels, out_channels):
        super(UNet, self).__init__()
        # Downsampling
        self.down_conv1 = ConvBlock(in_channels, 64)
        self.down_conv2 = ConvBlock(64, 128)
        self.down_conv3 = ConvBlock(128, 256)
        self.down_conv4 = ConvBlock(256, 512)

        # Upsampling
        self.up_block1 = UpBlock(512, 256, 256) # Input 512, middle 256, concatenate with 256 from skip connection
        self.up_block2 = UpBlock(256, 128, 128) # Input 256, middle 128, concatenate with 128 from skip connection
        self.up_block3 = UpBlock(128, 64, 64)   # Input 128, middle 64, concatenate with 64 from skip connection

        # Final layer
        self.final_conv = nn.Conv2d(64, out_channels, kernel_size=1)

```



```
def forward(self, x):  
    # Downsampling  
    x1 = self.down_conv1(x)  
    x2 = self.down_conv2(F.max_pool2d(x1, 2))  
    x3 = self.down_conv3(F.max_pool2d(x2, 2))  
    x4 = self.down_conv4(F.max_pool2d(x3, 2))  
  
    # Upsampling  
    x = self.up_block1(x4, x3)  
    x = self.up_block2(x, x2)  
    x = self.up_block3(x, x1)  
  
    # Final layer  
    out = self.final_conv(x)  
    return out
```

```
In [9]: n_channels = 4 # Number of channels in the input image (e.g., RGB+NIR)
```

```
# The number of output channels is 2 for binary classification  
model = UNet(n_channels, 2)  
  
loss_fn = nn.CrossEntropyLoss()
```

In [10]: `print(model)`

```

UNet(
  (down_conv1): ConvBlock(
    (conv1): Conv2d(4, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (down_conv2): ConvBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (down_conv3): ConvBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (down_conv4): ConvBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (up_block1): UpBlock(
    (up): ConvTranspose2d(512, 256, kernel_size=(2, 2), stride=(2, 2))
    (conv_block): ConvBlock(
      (conv1): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (up_block2): UpBlock(
    (up): ConvTranspose2d(256, 128, kernel_size=(2, 2), stride=(2, 2))

```

```
(conv_block): ConvBlock(
  (conv1): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(up_block3): UpBlock(
  (up): ConvTranspose2d(128, 64, kernel_size=(2, 2), stride=(2, 2))
  (conv_block): ConvBlock(
    (conv1): Conv2d(128, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
)
(final_conv): Conv2d(64, 2, kernel_size=(1, 1), stride=(1, 1))
)
```



```
In [11]: def train_model(model, train_dl, valid_dl, loss_fn, optimizer, epochs, n_classes):

    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    model.to(device)

    for epoch in range(epochs):
        print(f'Epoch {epoch}/{epochs - 1}')
        print('-' * 10)

        # Training phase
        model.train()
        train_loss, train_corrects, train_pixels = 0, 0, 0
        for inputs, labels in train_loader:
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model(inputs)

            # Debugging purpose
            if torch.any(labels >= n_classes):
                print(f"Invalid label detected. Labels should be less than {n_classes}.")

            optimizer.zero_grad()

            outputs = model(inputs)
            loss = loss_fn(outputs, labels)

            loss.backward()
            optimizer.step()

            # Convert the model's output to predicted classes
            preds = outputs.argmax(dim=1)

            # Update training metrics
            train_loss += loss.item() * inputs.size(0)
            train_corrects += torch.sum(preds == labels.data)
            train_pixels += torch.numel(labels)

        # Calculate and print training loss and accuracy
        train_epoch_loss = train_loss / len(train_dl.dataset)
        train_epoch_acc = train_corrects.double() / train_pixels
        print(f'Train Loss: {train_epoch_loss:.4f} Acc: {train_epoch_acc:.4f}')
```

```

# Validation phase
model.eval()
valid_loss, valid_corrects, valid_pixels = 0, 0, 0
with torch.no_grad():
    for inputs, labels in valid_dl:
        inputs, labels = inputs.to(device), labels.to(device)

        outputs = model(inputs)
        loss = loss_fn(outputs, labels)

        preds = outputs.argmax(dim=1)
        valid_loss += loss.item() * inputs.size(0)
        valid_corrects += torch.sum(preds == labels.data)
        valid_pixels += torch.numel(labels)

valid_epoch_loss = valid_loss / len(valid_dl.dataset)
valid_epoch_acc = valid_corrects.double() / valid_pixels
print(f'Valid Loss: {valid_epoch_loss:.4f} Acc: {valid_epoch_acc:.4f}')

print('Training complete')

```

In [12]: `optimizer = torch.optim.Adam(model.parameters(), lr=0.001)`

WARNING:tensorflow:From C:\Users\HP\anaconda3\Lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

WARNING:tensorflow:From C:\Users\HP\anaconda3\Lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

```
In [13]: epochs = 30 # Set the number of epochs
n_classes = 2
train_model(model, train_loader, val_loader, loss_fn, optimizer, epochs, n_classes)
```

Epoch 0/29

Train Loss: 0.3824 Acc: 0.8126

Valid Loss: 0.3080 Acc: 0.8727

Epoch 1/29

Train Loss: 0.3489 Acc: 0.8355

Valid Loss: 0.2606 Acc: 0.9022

Epoch 2/29

Train Loss: 0.3234 Acc: 0.8496

Valid Loss: 0.2405 Acc: 0.8951

Epoch 3/29

Train Loss: 0.2935 Acc: 0.8719

Valid Loss: 0.1792 Acc: 0.9333

Epoch 4/29

Train Loss: 0.2775 Acc: 0.8835

Valid Loss: 0.1680 Acc: 0.9405

Epoch 5/29

KeyboardInterrupt

Traceback (most recent call last)

Cell In[13], line 3

1 epochs = 30 # Set the number of epochs

2 n_classes = 2

----> 3 train_model(model, train_loader, val_loader, loss_fn, optimizer, epochs, n_classes)

Cell In[11], line 33, in train_model(model, train_dl, valid_dl, loss_fn, optimizer, epochs, n_classes)

30 preds = outputs.argmax(dim=1)

32 # Update training metrics

---> 33 train_loss += loss.item() * inputs.size(0)

34 train_corrects += torch.sum(preds == labels.data)

35 train_pixels += torch.numel(labels)

KeyboardInterrupt:

In []: